

# Introduction to Systems Programming

COMP2100

Week 1

Richard Han/**Young** Choon Lee

4 Research Park Drive

[richard.han@mq.edu.au](mailto:richard.han@mq.edu.au)/[young.lee@mq.edu.au](mailto:young.lee@mq.edu.au)

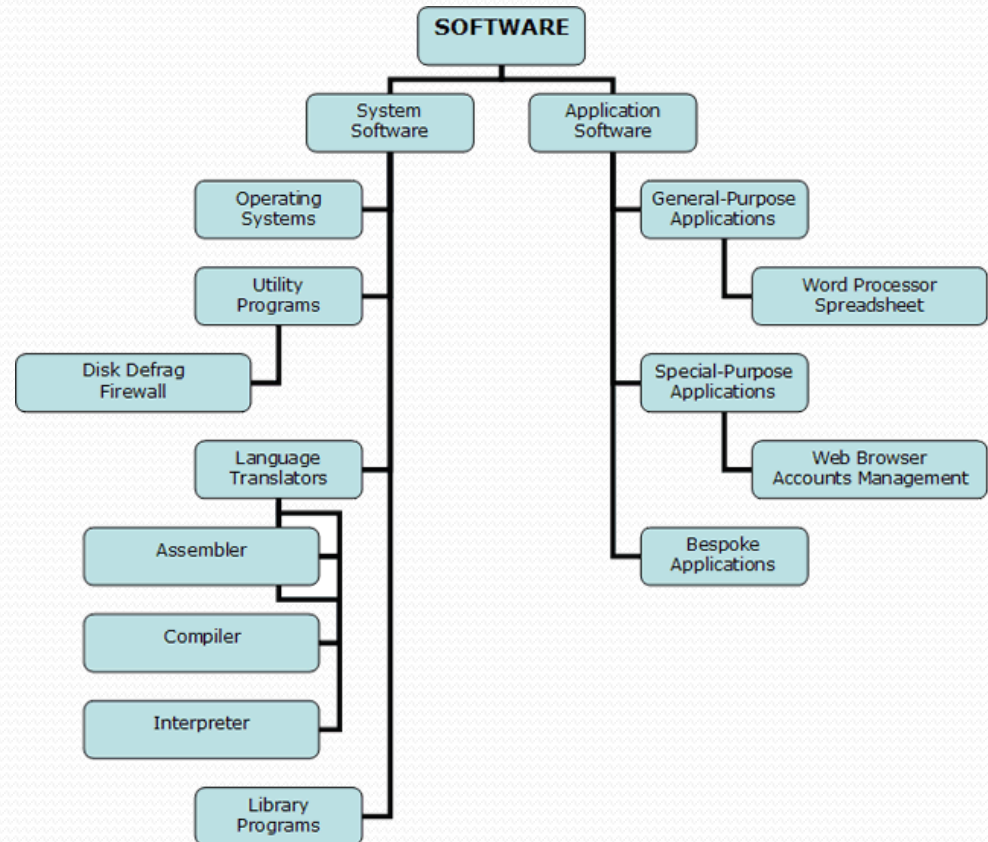

# Systems programming **vs** Application programming



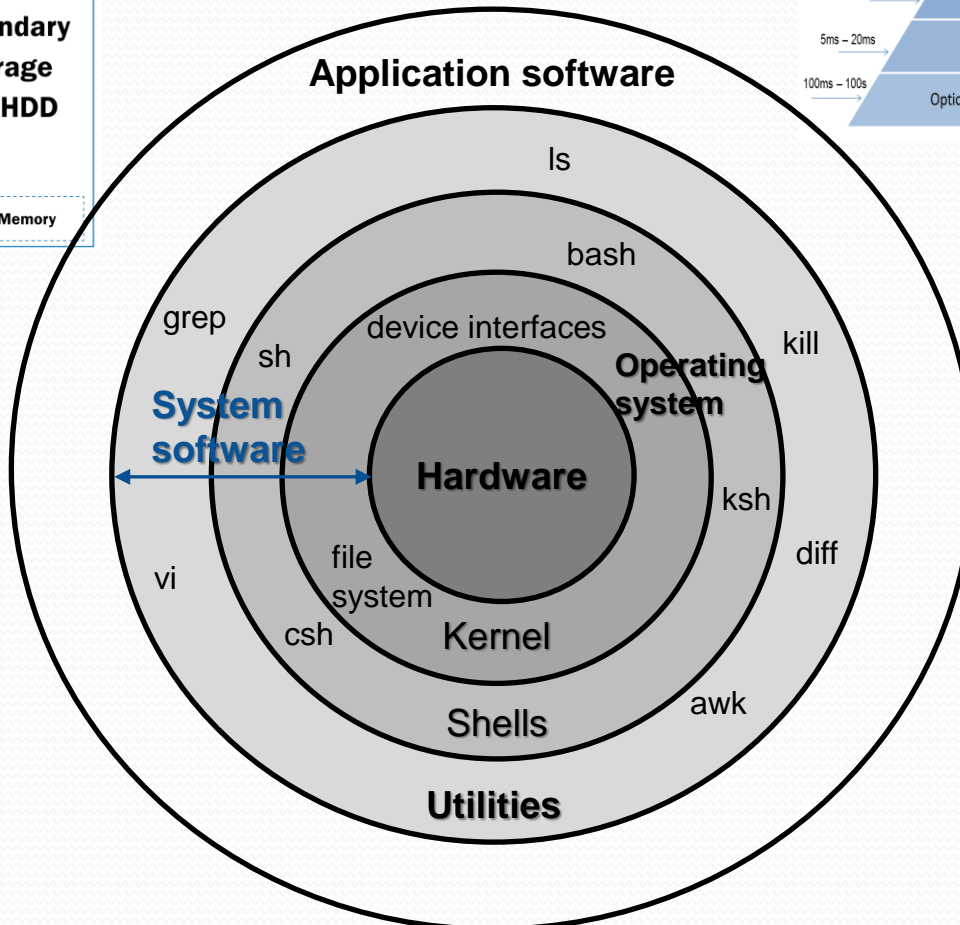

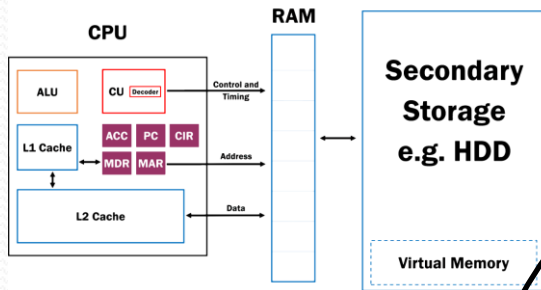
Application programming:  
facing (end-)users

Systems programming:  
facing systems

# System software



# Know the system



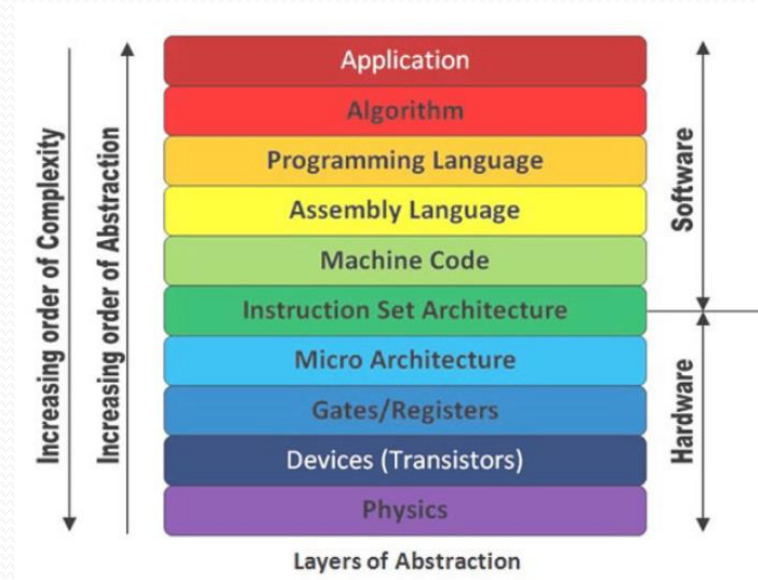
```

00000000 01 00 FF FF 00 00 00 00 00 00 00 00 40 00 CC 00 .....@S
00000010 0C 00 00 00 00 00 00 26 01 8F 00 00 00 00 00 53 00 .....elect NS
00000020 65 00 6C 00 65 00 63 00 74 00 20 00 52 00 75 00 .....le Shell D
00000030 6C 00 65 00 00 00 00 00 00 00 00 00 01 4D 00 53 00 .....lv
00000040 20 00 53 00 68 00 65 00 6C 00 6C 00 20 00 44 00 .....PS : 6.2%
00000050 6C 00 67 00 00 00 00 00 00 00 00 00 00 02 00 00 .....&Apply
00000060 03 01 A1 50 53 00 3A 00 C3 00 26 00 32 25 00 00 .....F VA 36
00000070 FF FF 63 00 00 00 00 00 00 00 00 00 00 00 00 .....&Apply
00000080 03 00 01 50 00 00 54 00 41 00 0A 00 4A 26 00 00 .....to all P
00000090 FF FF 00 00 74 00 41 00 70 00 70 00 6C 00 79 00 .....O.K.
000000A0 20 00 74 00 6F 00 20 00 61 00 6C 00 6C 00 00 00 .....F.2
000000B0 00 03 00 00 00 00 00 00 00 00 00 00 61 00 61 50 .....Cancel P
000000C0 4F 00 4E 00 00 00 00 00 00 00 00 00 FF FF 80 00 .....
000000D0 7E 00 7D 00 32 00 0E 00 01 00 00 00 FF FF 80 00 .....
000000E0 00 00 01 50 04 00 7D 00 32 00 0E 00 02 00 00 00 .....F.2
000000F0 FF FF 00 00 43 00 61 00 6E 00 63 00 65 00 6C 00 .....Cancel P
00000100 00 00 00 00 00 00 00 00 00 00 00 00 00 00 01 50 .....
00000110 EA 00 7D 00 32 00 0E 00 09 00 00 00 FF FF 80 00 .....&Help
00000120 24 00 48 00 65 00 6C 00 70 00 00 00 00 00 00 00 .....%
00000130 00 00 00 00 00 00 00 00 00 00 81 50 0E 00 3A 00 .....F
00000140 00 00 0E 00 2F 75 00 00 FF FF 81 00 00 00 00 00 .....%
00000150 00 00 00 00 00 00 00 00 00 00 62 50 0E 00 30 00 .....F.0
00000160 1E 00 00 00 8E 75 00 00 FF FF 82 00 46 00 69 00 .....ie Type P
00000170 6C 00 65 00 20 00 54 00 79 00 70 00 65 00 00 00 .....T.O.
00000180 00 00 00 00 00 00 00 00 00 00 00 00 00 00 02 50 .....Parsing
00000190 54 00 30 00 2C 00 00 00 2F 25 00 00 FF FF 82 00 .....Rules
000001A0 50 00 61 00 72 00 73 00 69 00 4E 00 67 00 20 00 .....q %
000001B0 52 00 75 00 6C 00 65 00 73 00 00 00 00 00 00 00 .....F
000001C0 00 00 00 00 00 00 00 00 07 00 00 50 06 00 07 00 .....q %
000001D0 1A 01 71 00 8D 75 00 00 FF FF 80 00 00 00 00 00 .....F
000001E0 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....F
000001F0 3E 00 00 00 8C 75 00 00 FF FF 82 00 53 00 65 00 .....F
00000200 6C 00 65 00 63 00 74 00 20 00 52 00 75 00 6C 00 .....lect Rul
00000210 65 00 20 00 46 00 6F 00 72 00 20 00 46 00 69 00 .....e For Fi
00000220 6C 00 65 00 00 00 00 00 00 00 00 00 00 00 00 00 .....le
00000230 00 00 01 50 0E 00 1B 00 00 01 0E 00 20 25 00 00 .....F.0
00000240 FF FF 01 00 00 00 00 00 00 00 00 00 00 00 00 00 .....F.a? m
00000250 00 00 02 50 14 00 61 00 27 00 00 00 48 26 00 00 .....
00000260 FF FF 92 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
    
```

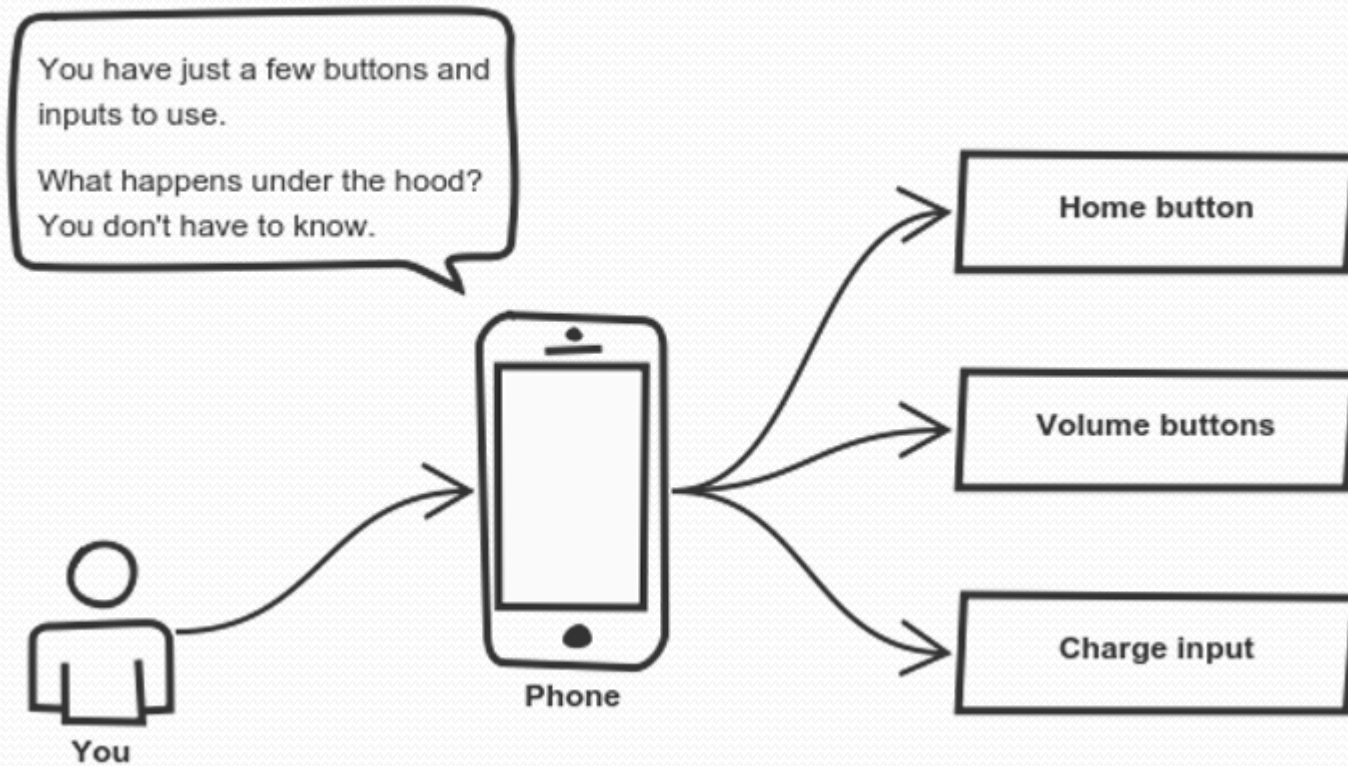
# Abstraction

- Simplification of complex system via interface (modelling)
  - Machine language
  - Assembly: an abstraction of the machine language
  - Many languages are abstraction of assembly language, e.g., Java, Python and C


*\* But they still require you to think in terms of the **structure of the computer** rather than the **structure of the problem***



# Abstraction



# Reality vs Abstraction

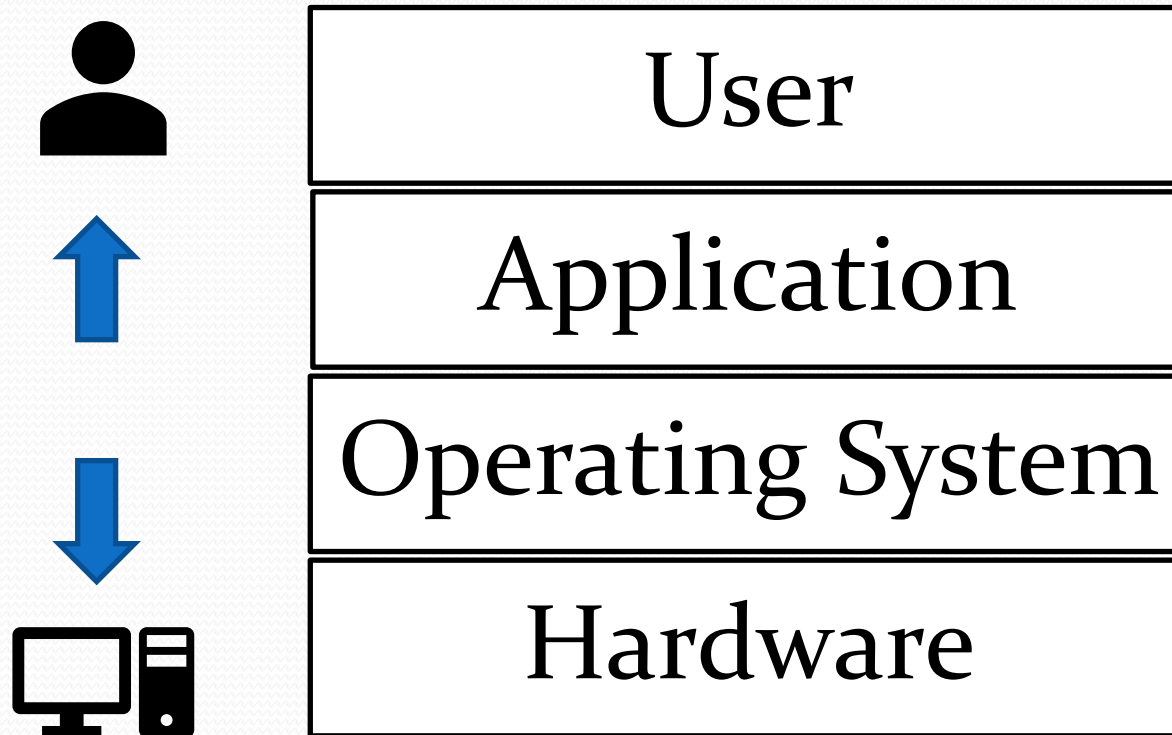






# Systems programming and C programming language



- C is a native language for many operating systems



# Software Systems written in C and the like

- Operating systems
  - Unix/Linux-like: mostly C and C++
  - Windows: C, C++, C# and some assembly code
  - OSX/iOS: C, C++, Objective-C
  - Android: C, C++, Java
- Other application software
  - MS Office: C → C++
  - Adobe Photoshop: C++
  - Web browsers inc. Google Chrome, Firefox, Opera: C++

# C: old == obsolete?

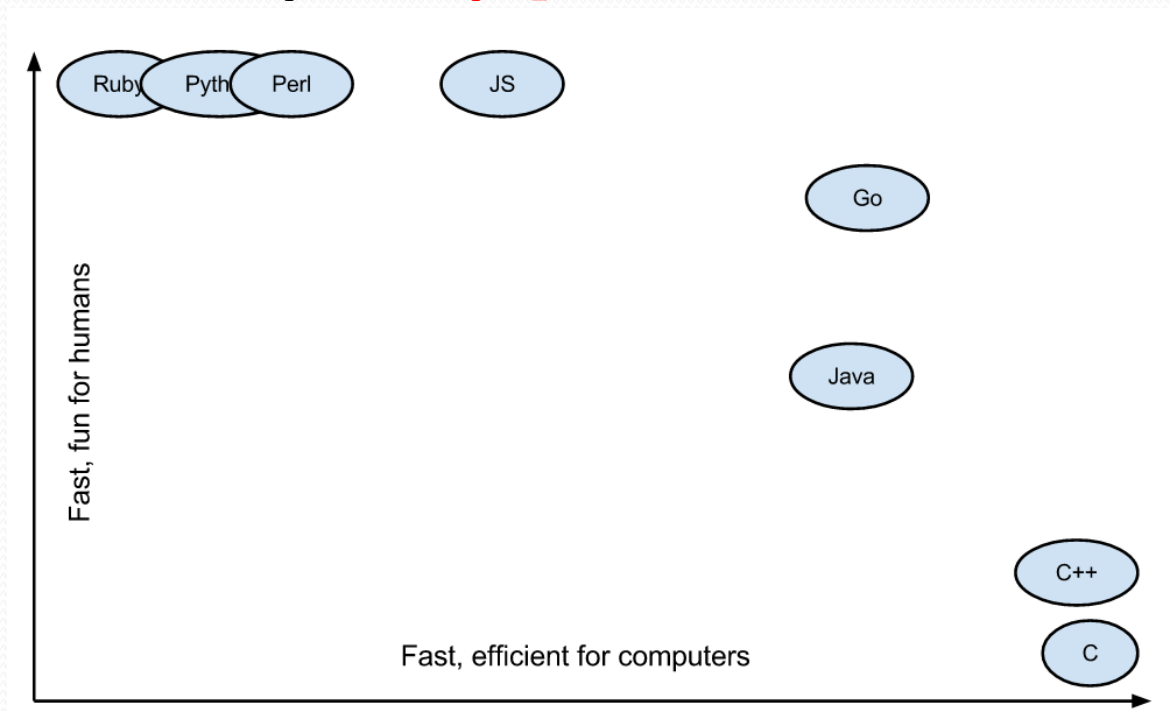


# C for Java programmers

- Java is mid-90s high-level OO language
- C is early-70s *procedural* language
- C advantages:
  - Direct access to OS primitives (system calls)
  - Fewer library issues – just execute
- (More) C disadvantages:
  - language is portable, APIs are not (e.g., `unistd.h`)
  - memory and “handle” leaks
  - preprocessor can lead to obscure errors

# Java vs C

“Although C is often criticized for its accidental complexity, unsafe programming, and lack of features. Also, C is platform-dependent, i.e., C code is not portable. But **if you want to make the most use of your hardware, then C/C++ or Rust is your **only option**.**” \*



\* <https://towardsdatascience.com/top-10-in-demand-programming-languages-to-learn-in-2020-4462eb7d8d3e>

# Operators in C: C == Java?

- Arithmetic      + - \* /    unary -
- Bit-wise        ~ & | ^    << >>
- Relational      < <= == != >= >
- Logical         ! && ||

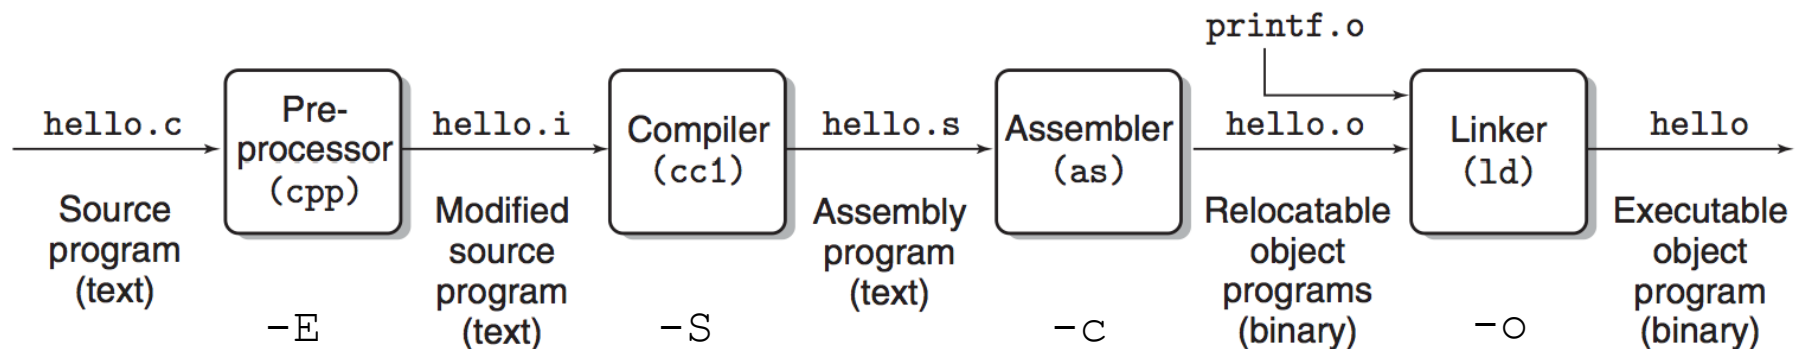
# C vs Java

```
public class hello
{
    public static void main (String args []) {
        System.out.println("Hello world");
    }
}
```

```
#include <stdio.h>
int main(int argc, char *argv[])
{
    printf("Hello, World\n");
    return 0;
}
```

# Executing C programs

- Java programs semi-interpreted:
  - `javac` converts `.java` into `.class`
  - not machine-specific
  - *byte codes* are then interpreted by JVM
- C programs are normally compiled and linked:
  - `gcc` converts `.c` into `a.out`
  - `a.out` is executed by OS and hardware

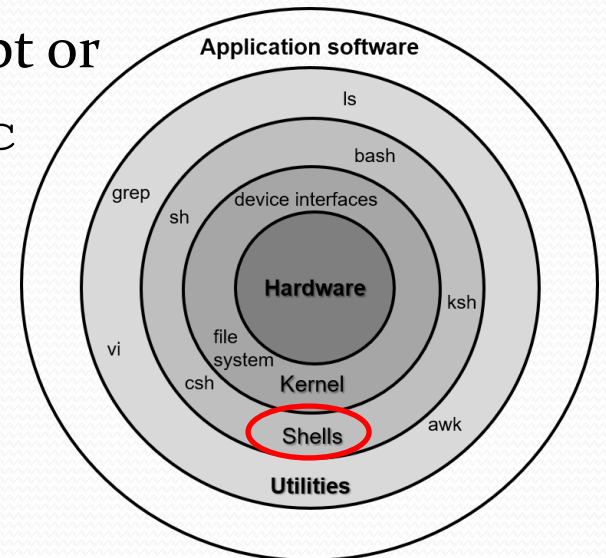






# Command Line Interface (CLI)

- A (UNIX) command is a program which interacts with the OS kernel to provide the environment and perform the functions called for by the user.
- A command can be:
  - a built-in shell command, e.g., `ls`, `cp` and `whoami`
  - an executable shell file, aka a shell script or
  - a source compiled, object code file, `gcc`
- The shell is a command line interpreter



# CLI: Man Pages

- The **man** command allows you to access the MANual pages for a UNIX command.
- To get additional help on any of the commands listed below, you can always type **man name\_of\_command** at the command prompt.
- Examples:
  - **man ssh**
  - **man passwd**

# CLI: Commands

- **ls** : lists the contents of a directory
  - **l** : long directory listing
  - **a** : lists all files, including files which are normally hidden
  - **F** : distinguishes between directories and regular files
  - **h** : ? Look it up using **man**
- **pwd** : prints the current working directory
- **cd** : changes directories
  - The difference between relative and absolute paths.
  - Special characters **.**, **..**, and **~**.
- **mkdir** : creates a directory
- **rmdir** : removes a directory (assuming it is empty)
  - If you get an error that the directory isn't empty even though it looks empty, check for hidden files.

# CLI: Commands

- **touch** : creates an empty file with the specified name, or if the file already exists it modifies the timestamp.
- **rm** : removes a file.
  - f : force deletion
  - r : recursive deletion
- **mv** - moves a file, or renames a file
  - f : forces overwrite, if the destination file exists
- **cp** - copies a file, leaving the original intact
  - f : forces overwrite, if the destination file exists
  - r : recursive copying of directories

# CLI: The UNIX Pipe (|)

- The pipe (|) creates a channel from one command to another. Think of the pipe as a way of connecting the output from one command to the input of another command.
- The pipe can be used to link commands together to perform more complex tasks that would otherwise take multiple steps (and possibly writing information to disk).
- Example: count the number of users logged onto the current system.
  - `who | wc -l`

# CLI: Redirection (<, >, >>)

- <: redirects input
  - `command < file`
  - `command` takes its input from `file`
- >: redirects output (overwrite)
  - `command > file`
  - `command` (over)writes its output from `file`
- >>: redirects output (append)
  - `command >> file`
  - `command` writes/appends its output from `file`

# CLI: Quotes

- ' (single quote)
  - Enclosing characters in single quotes (') preserves the literal value of each character within the quotes.
- " (double quote)
  - Enclosing characters in double quotes (") preserves the literal value of all characters within the quotes, with the exception of \$, `, \, and, when history expansion is enabled, !.

*\* Anything which is enclosed between single and double quote can be treated as single string.*
- ` (back quote)
  - Command substitution
- \ (back slash)
  - Backslash tells shell to "ignore next character."



# CLI: Quotes

```
unit=comp2100
```

#	Expression	Result	Comments
1	"\$unit"	comp2100	variables are expanded inside ""
2	`\$unit`	\$unit	variables are not expanded inside ` `
3	""'\$unit'""	`comp2100`	' ' has no special meaning inside ""
4	```\$unit```	"\$unit"	"" is treated literally inside ``
5	`\"`	`\"`	\ has no special meaning inside ``
6	""\'""	`\'`	`\'` is interpreted inside "" but has no significance for `
7	""\"""	\"	\" is interpreted inside ""
8	""*""	*	glob does not work inside "" or ``
9	"`echo hi`"	hi	`` and \$() are evaluated inside ""
10	`echo hi`	`echo hi`	`` and \$() are not evaluated inside ``
11	`!gcc`	!gcc	history expansion character `!` is ignored inside ``
12	"!gcc"	gcc args	expands to the most recent command matching "cmd"



# Basic C Data Types

Type Name	Size in bits	Literal	Literal
char	8	'A'	65
short int	16	25000	-25000
int	32	2000000000	-2000000000
long int	64	900000000000	-9000000000000000
unsigned char	8	255	0xff
unsigned short	16	50000	0xface
unsigned int	32	4000000000	0x12345678
unsigned long	64	900000000000	0x1234567812345678
float	32	1.34f	
double	64	3.1415926535	1.234e+5
<i>Pointer</i>	64	N/A	

# Number systems

<u>Binary Digit</u>	<u>Octal Digit</u>	<u>Decimal Digit</u>	<u>Hexadecimal Digit</u>
0000	0	0	0
0001	1	1	1
0010	2	2	2
0011	3	3	3
0100	4	4	4
0101	5	5	5
0110	6	6	6
0111	7	7	7
1000	10	8	8
1001	11	9	9
1010	12	10	A
1011	13	11	B
1100	14	12	C
1101	15	13	D
1110	16	14	E
1111	17	15	F

# Data types: Java vs C

Type	Java	C
char	16 bits	8 bits
short	16 bits	16 bits
int	32 bits	16, 32 or 64 bits
long	64 bits	32 or 64 bits
float	32 bits	32 bits
double	64 bits	64 bits
boolean	1 bit	(use int)
byte	8 bits	(use char)
long long		64 bits (unofficial)
long double		80, 96 or 128 bits

# Data in C

- Arrays: no automatic initialisation in C

- C:** `int array[5];`

??	??	??	??	??
----	----	----	----	----

`sizeof()`

- Java:** `int[] array = new int[5];`

5	00	00	00	00	00
---	----	----	----	----	----

`array.length`

- Characters and strings: string bounded by a `'\0'` (null)

- Example: the string (char array) "COMP2100"

- C (ASCII)

43	4F	4D	50	32	31	30	30	\0
----	----	----	----	----	----	----	----	----

- Java (Unicode)

8	00	43	00	4F	00	4D	00	50	00	32	00	31	00	30	00	30
---	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

- struct vs object

# Formatted Output: printf

```
int printf(char *format, arg1, arg2, ...);
```

- `printf` converts, formats, and prints its arguments on the standard output under control of the `format`. It returns the number of characters printed.
- `format` contains
  - Ordinary characters
  - Conversion specifications
    - Syntax: `%[Flag][Width][Prec][Length]Specifier`
    - e.g., `%c`, `%d`, `%s`, `%x`

# Basic conversion specifications

Character (format specifier)	Argument type; printed as
d, i	int; decimal number
o	int; unsigned octal number (without a leading zero)
x, X	int; unsigned hexadecimal number (without a leading ox or oX), using abcdef or ABCDEF for 10, ...,15.
u	int; unsigned decimal number
c	int; single character
s	char *; print characters from the string until a '\0' or the number of characters given by the precision.
f	double; [-]m.dddddd, where the number of d's is given by the precision (default 6).
e, E	double; [-]m.dddddde+/-xx or [-]m.dddddde+/-xx, where the number of d's is given by the precision (default 6).
g, G	double; use %e or %E if the exponent is less than -4 or greater than or equal to the precision; otherwise use %f. Trailing zeros and a trailing decimal point are not printed.
p	void *; pointer (implementation-dependent representation).
%	no argument is converted; print a %



# printf: examples

```
#include <stdio.h>
int main()
{
    char ch = '@';
    char str[30] = "COMP2100 Systems Programming";
    float pi = 3.14;
    int unit = 2100;
    double lpi = 3.141592;
    printf("Character is %c \n", ch);
    printf("String is %s \n" , str);
    printf("Float value is %f \n", pi);
    printf("Integer value is %d\n" , unit);
    printf("Double value is %lf \n", lpi);
    printf("Octal value is %o \n", unit);
    printf("Hexadecimal value is %x \n", unit);
    return 0;
}
```



# Functions

Prototype

```
int sum(int a, int b);
```

```
int main(int argc, char **argv) {
```

```
...
```

```
    j = sum(3, 5);
```

```
...
```

```
}
```

```
int sum(int x, int y) {
```

```
    return x + y;
```

```
}
```

Definition